

# Assignment 2: Classification for the Detection of Opinion Spam

Anouk van der Lee (6620590), Shu Zhao (6833519), Fleur Petit (5583837)

12, October, 2020

## Contents

<b>TO DO</b>	<b>2</b>
<b>0. Pre-processing</b>	<b>2</b>
Load reviews . . . . .	2
Prepare corpus . . . . .	3
Get word frequency matrix . . . . .	3
Vector with labels . . . . .	4
Random split in train and test data . . . . .	4
<b>1. Multinomial naive Bayes (generative linear classifier)</b>	<b>4</b>
Function for multinomial Bayes classifier . . . . .	4
Train . . . . .	5
Predict . . . . .	5
Confusion matrix . . . . .	6
Performance metrics . . . . .	6
<b>2. Regularised logistic regression (discriminative linear classifier)</b>	<b>7</b>
Choose lambda . . . . .	7
Train . . . . .	8
Predict . . . . .	8
Confusion matrix . . . . .	8
Performance metrics . . . . .	9
<b>3. Classification trees (flexible classifier)</b>	<b>9</b>
Train . . . . .	9
Predict . . . . .	9
Confusion matrix . . . . .	9
Performance metrics . . . . .	10
<b>4. Random forests (flexible classifier)</b>	<b>10</b>
Predict . . . . .	10
Confusion matrix . . . . .	10
Performance metrics . . . . .	11
<b>Comparison</b>	<b>11</b>
Performance metrics . . . . .	11
Logistic regression . . . . .	12

## TO DO

- Normalise data before training
  - Check if necessary: <https://stats.stackexchange.com/questions/29781/when-conducting-multiple-regression-when-should-you-center-your-predictor-variables>
- Cross-validate/choose random forest parameters
  - ntree
  - mtry
- Visualise more of the results
- Write more explanations, figure captions, table captions, etcetera
  - figure captions are done in the chunk-header
  - kable captions are done in the 'kable()' function.
  - Latex code can be used for references and labels
    - \* for the tables: <https://stackoverflow.com/questions/54082814/adding-label-in-kable-kableextra-latex-output>
    - \* for the figures: <https://cran.r-project.org/web/packages/office2vignettes/vignettes/captions.html>

## 0. Pre-processing

### Load reviews

```
df <- NULL

file_names <-
  list.files(path = "negative_polarity", recursive = T)

for (name in file_names)
  df <- bind_rows(df,
    list(label = as.numeric(str_detect(name, "truthful")), # 0 for deceptive and 1 for truthful
          fold = as.numeric(str_extract(name, pattern = regex("\\d+"))),
          hotel = str_extract(name, pattern = regex("[a-z]*_\\d+")),
          review = read_file(str_c("negative_polarity/", name))
    )
  )

df <- df %>%
  mutate(id = as.numeric(str_extract(hotel, pattern = regex("\\d+"))),
         hotel = str_extract(hotel, pattern = regex("[a-z]*"))
  )

head(df, 2) %>%
  kable()
```

label	document	id
0	<p>1 hilton We stayed at the Schicago Hilton for 4 days and 3 nights for a conference. I have to say, normally I am very easy going about amenities, cleanliness, and the like. . . however our experience at the Hilton was so awful I am taking the time to actually write this review. Truly, DO NOT stay at this hotel. When we arrived in our room, it was clear that the carpet hadn't been vacuumed. I figured, "okay, it's just the carpet." Until I saw the bathroom! Although the bathroom had all the superficial indicators of housekeeping having recently cleaned (i.e., a paper band across the toilet, paper caps on the drinking glasses, etc., it was clear that no ACTUAL cleaning took place. There was a spot (probably urine!) on the toilet seat and, I kid you not, the remnants of a lip-smudge on the glass. I know people who have worked many years in the hotel industry and they always warned that lazy housekeeping will make things "appear" clean but in fact they make no effort to keep things sanitary. Well, the Hilton was proof. I called downstairs and complained, and they sent up a chambermaid hours later. Frankly, I found the room disgusting. The hotel itself, outside the rooms, was cavernous and unwelcoming, with an awful echo in the lobby area that created a migraine-inducing din. Rarely have I been so eager to leave a place as this. When I got home, I washed all my clothes whether I had worn them or not, such was the skeeviness of our accomodations. Please, do yourself a favor and stay at a CLEAN hotel.</p> <p>1 hilton Hotel is located 1/2 mile from the train station which is quite hike when you're traveling with luggage and/or kids. They seem to cash in on guests who arrive in private car by charging exorbitant parking/valet fees. Rooms feature either double or king sized beds; no queen beds at all. If you want a little extra leg room in your bed, the price jump from double- to king-sized is stiff. Rooms with any kind of view pay a healthy surcharge, too.</p>	1

## Prepare corpus

```
corpus <- VCorpus(VectorSource(df[["review"]])) %>%
  tm_map(., content_transformer(tolower)) %>% # no capital letters
  tm_map(., stripWhitespace) %>% # remove extra white space
  tm_map(., removeWords, stopwords("english")) %>% # remove stopwords
  tm_map(., stemDocument) # stem words
```

## Get word frequency matrix

```
dtm <- DocumentTermMatrix(corpus) %>%
  removeSparseTerms(., sparse = 0.7) # those terms from x are removed which have at least a 70
# percent of empty (i.e., terms occurring 0 times in a document) elements. I.e., the resulting
# matrix contains only terms with a sparse factor of less than 70 percent

inspect(dtm)
```

```
## <<DocumentTermMatrix (documents: 800, terms: 10)>>
## Non-/sparse entries: 3754/4246
## Sparsity : 53%
## Maximal term length: 7
## Weighting : term frequency (tf)
## Sample :
## Terms
## Docs  chicago get hotel like one room servic staff stay will
## 128      2  3      3  1  2  6      1  2  1  0
## 18       0  3      3  3  1  5      0  2  6  2
```

```
## 247      1  0      9  4  2  3      1  0  0  1
## 311      1  2      5  0  3  4      1  0  6  0
## 390      1  1      4  0  5  6      0  0  4  3
## 412      3  0      2  0  4 11      0  0  2  1
## 514      0  4      4  2  4  4      0  2  1  1
## 549      3  4      4  5  1  6      4  0  3  0
## 610      0  5      8  0  3  9      3  0  4  1
## 7        1  5      4  3  3  5      0  0  1  0
```

## Vector with labels

```
labels <- df[["label"]]

labels[c(1:10, 790:800)]

## [1] 0 0 0 0 0 0 0 0 0 0 1 1 1 1 1 1 1 1 1 1
```

## Random split in train and test data

```
train_ind <- sample(1:nrow(df), size = nrow(df)*0.75)

train_dtm <- as.matrix(dtm)[train_ind,]
train_labels <- labels[train_ind]

test_dtm <- as.matrix(dtm)[-train_ind,]
test_labels <- labels[-train_ind]
```

# 1. Multinomial naive Bayes (generative linear classifier)

## Function for multinomial Bayes classifier

```
# Function for multinomial naive Bayes

train.mnb <- function (dtm, labels) {
  call <- match.call()
  V <- ncol(dtm)
  N <- nrow(dtm)
  prior <- table(labels) / N
  labelnames <- names(prior)
  nclass <- length(prior)
  cond.probs <- matrix(nrow = V, ncol = nclass)
  dimnames(cond.probs)[[1]] <- dimnames(dtm)[[2]]
  dimnames(cond.probs)[[2]] <- labelnames
  index <- list(length = nclass)
  for (j in 1:nclass) {
    index[[j]] <- c(1:N)[labels == labelnames[j]]
  }

  for (i in 1:V) {
    for (j in 1:nclass) {
      cond.probs[i, j] <-
        (sum(dtm[index[[j]], i]) + 1) / (sum(dtm[index[[j]], ]) + V)
    }
  }
}
```

```

}
list(call = call,
     prior = prior,
     cond.probs = cond.probs)
}

predict.mnb <- function (model, dtm) {
  classlabels <- dimnames(model$cond.probs)[[2]]
  logprobs <- dtm %*% log(model$cond.probs)
  N <- nrow(dtm)
  nclass <- ncol(model$cond.probs)
  logprobs <-
    logprobs + matrix(nrow = N,
                      ncol = nclass,
                      log(model$prior),
                      byrow = T)
  classlabels[max.col(logprobs)]
}

```

## Train

```

mnb_model <- train.mnb(train_dtm, train_labels)

mnb_model

## $call
## train.mnb(dtm = train_dtm, labels = train_labels)
##
## $prior
## labels
##      0      1
## 0.4783333 0.5216667
##
## $cond.probs
##      0      1
## chicago 0.08410867 0.03902644
## get      0.06140677 0.07427612
## hotel    0.22590249 0.21947125
## like     0.05768515 0.04825850
## one      0.05582434 0.06924045
## room     0.22925195 0.25891733
## servic   0.05694083 0.05874948
## staff    0.04205434 0.04616030
## stay     0.14551544 0.14141838
## will     0.04131001 0.04448175

```

## Predict

```

predicted_mnb <- predict.mnb(mnb_model, test_dtm)

```

## Confusion matrix

```
conf_mnb <- table(test_labels, predicted_mnb)
```

```
conf_mnb
```

```
##           predicted_mnb
## test_labels  0  1
##           0 55 58
##           1 18 69
```

## Performance metrics

```
performance <- function(confusion_matrix){
  tn <- confusion_matrix[1,1]
  fn <- confusion_matrix[2,1]
  n_pred <- tn+fn

  fp <- confusion_matrix[1,2]
  tp <- confusion_matrix[2,2]
  p_pred <- fp+tp

  n <- tn+fp
  p <- fn+tp

  performance_metrics <- tibble(metric = c("recall",
                                           "miss-rate", # 1 - recall
                                           "fall-out", # 1 - selectivity
                                           "selectivity",
                                           "prevalence",
                                           "precision",
                                           "false omission rate", # 1 - neg_predictive_value
                                           "pos likelihood ratio",
                                           "neg likelihood ratio",
                                           "accuracy",
                                           "false discovery rate", # 1 - precision
                                           "neg predictive value",
                                           "diagnostic odds ratio",
                                           "F1"),
                                value = c(tp/p, # recall
                                           fn/p, # 1 - recall
                                           fp/n, # 1 - selectivity
                                           tn/n, # selectivity
                                           p/(n+p), # prevalence
                                           tp/p_pred, # precision
                                           fn/n_pred, # 1 - neg_predictive_value
                                           (tp/p)/(fp/n), # positive likelihood ratio
                                           (fn/p)/(tn/n), # negative likelihood ratio
                                           (tp+tn)/(n+p), # accuracy
                                           fp/p_pred, # 1 - precision
                                           tn/n_pred, # negative predictive value
                                           ((tp/p)/(fp/n))/((fn/p)/(tn/n)), # diagnostic odds ratio
                                           2*(tp/p_pred)*(tp/p)/((tp/p_pred) + (tp/p))) # F1
  )
}
```

```

  performance_metrics
}

perf_mnb <-performance(conf_mnb)

perf_mnb %>% kable()

```

metric	value
recall	0.7931034
miss-rate	0.2068966
fall-out	0.5132743
selectivity	0.4867257
prevalence	0.4350000
precision	0.5433071
false omission rate	0.2465753
pos likelihood ratio	1.5451843
neg likelihood ratio	0.4250784
accuracy	0.6200000
false discovery rate	0.4566929
neg predictive value	0.7534247
diagnostic odds ratio	3.6350575
F1	0.6448598

## 2. Regularised logistic regression (discriminative linear classifier)

Choose lambda

```


# Use LASSO (penalising for number of parameters)
# Determine lambda by means of cross validation

set.seed(1)

cv_lasso <- cv.glmnet(train_dtm, train_labels, alpha = 1, family = "binomial")

plot(cv_lasso)

```



## Train

```
model_glm <- glmnet(train_dtm, train_labels, alpha = 1, family = "binomial",
  lambda = cv_lasso$lambda.1se) # I choose the largest lambda within 1se from the sma
```

## Predict

```
#test_dtm_glm <- model.matrix(test_labels ~ test_dtm)[-1]

probabilities_glm <- predict(model_glm,
  newx = test_dtm,
  type = "response", # Type "response" gives the fitted probabilities for "b
  s = cv_lasso$lambda.1se)

predicted_glm <- ifelse(probabilities_glm > 0.5, 1, 0)
```

## Confusion matrix

```
conf_glm <- table(test_labels, predicted_glm)

conf_glm
```

```
##           predicted_glm
## test_labels  0  1
```



```
##          0  59  54
##          1  22  65
```

## Performance metrics

```
perf_glm <- performance(conf_glm)

perf_glm %>% kable()
```

metric	value
recall	0.7471264
miss-rate	0.2528736
fall-out	0.4778761
selectivity	0.5221239
prevalence	0.4350000
precision	0.5462185
false omission rate	0.2716049
pos likelihood ratio	1.5634312
neg likelihood ratio	0.4843172
accuracy	0.6200000
false discovery rate	0.4537815
neg predictive value	0.7283951
diagnostic odds ratio	3.2281145
F1	0.6310680

## 3. Classification trees (flexible classifier)

### Train

```
model_tree <- randomForest(x = train_dtm, y = train_labels, ntree = 1, mtry = ncol(train_dtm))

## Warning in randomForest.default(x = train_dtm, y = train_labels, ntree = 1, :
## The response has five or fewer unique values. Are you sure you want to do
## regression?
```

### Predict

```
probabilities_tree <- predict(model_tree, test_dtm, type = "response")

predicted_tree <- ifelse(probabilities_tree > 0.5, 1, 0)
```

### Confusion matrix

```
conf_tree <- table(test_labels, predicted_tree)

conf_tree
```

```
##          predicted_tree
## test_labels  0  1
##          0  67  46
##          1  36  51
```

## Performance metrics

```
perf_tree <- performance(conf_tree)

perf_tree %>% kable()
```

metric	value
recall	0.5862069
miss-rate	0.4137931
fall-out	0.4070796
selectivity	0.5929204
prevalence	0.4350000
precision	0.5257732
false omission rate	0.3495146
pos likelihood ratio	1.4400300
neg likelihood ratio	0.6978899
accuracy	0.5900000
false discovery rate	0.4742268
neg predictive value	0.6504854
diagnostic odds ratio	2.0634058
F1	0.5543478

## 4. Random forests (flexible classifier)

```
# TO DO: cross validate for mtry (caret package)

model_random <- randomForest(x = train_dtm, y = train_labels, ntree = 10, mtry = 5)

## Warning in randomForest.default(x = train_dtm, y = train_labels, ntree = 10, :
## The response has five or fewer unique values. Are you sure you want to do
## regression?
```

### Predict

```
probabilities_random <- predict(model_random, test_dtm, type = "response")

predicted_random <- ifelse(probabilities_random > 0.5, 1, 0)
```

### Confusion matrix

```
conf_random <- table(test_labels, predicted_random)

conf_random
```

```
##           predicted_random
## test_labels 0 1
##           0 66 47
##           1 32 55
```

## Performance metrics

```
perf_random <- performance(conf_random)

perf_random %>% kable()
```

metric	value
recall	0.6321839
miss-rate	0.3678161
fall-out	0.4159292
selectivity	0.5840708
prevalence	0.4350000
precision	0.5392157
false omission rate	0.3265306
pos likelihood ratio	1.5199315
neg likelihood ratio	0.6297457
accuracy	0.6050000
false discovery rate	0.4607843
neg predictive value	0.6734694
diagnostic odds ratio	2.4135638
F1	0.5820106

## Comparison

### Performance metrics

```
perf_compare <- tibble(metric = perf_mnb[["metric"]],
  mnb = perf_mnb[["value"]],
  glm = perf_glm[["value"]],
  tree = perf_tree[["value"]],
  random = perf_random[["value"]])

perf_compare %>% kable()
```

metric	mnb	glm	tree	random
recall	0.7931034	0.7471264	0.5862069	0.6321839
miss-rate	0.2068966	0.2528736	0.4137931	0.3678161
fall-out	0.5132743	0.4778761	0.4070796	0.4159292
selectivity	0.4867257	0.5221239	0.5929204	0.5840708
prevalence	0.4350000	0.4350000	0.4350000	0.4350000
precision	0.5433071	0.5462185	0.5257732	0.5392157
false omission rate	0.2465753	0.2716049	0.3495146	0.3265306
pos likelihood ratio	1.5451843	1.5634312	1.4400300	1.5199315
neg likelihood ratio	0.4250784	0.4843172	0.6978899	0.6297457
accuracy	0.6200000	0.6200000	0.5900000	0.6050000
false discovery rate	0.4566929	0.4537815	0.4742268	0.4607843
neg predictive value	0.7534247	0.7283951	0.6504854	0.6734694
diagnostic odds ratio	3.6350575	3.2281145	2.0634058	2.4135638
F1	0.6448598	0.6310680	0.5543478	0.5820106

## Logistic regression

```
predictions_per_model <- tibble(predictions = c(predicted_mnb,
                                              as.vector(predicted_glm),
                                              predicted_tree,
                                              predicted_random),
                                models = c(rep("mnb",
                                              length(predicted_mnb)),
                                           rep("logistic regression",
                                              nrow(predicted_glm)),
                                           rep("single tree",
                                              length(predicted_tree)),
                                           rep("random forest",
                                              length(predicted_random))),
                                ground_truth = c(rep(test_labels, 4)),
                                correct = ifelse(ground_truth == predictions, 1, 0)) %>%
mutate(models = factor(models, levels = c("mnb", "logistic regression", "single tree", "random forest")))

glm(correct ~ models, family = "binomial", data = predictions_per_model) %>%
  tidy() %>%
  kable()
```

term	estimate	std.error	statistic	p.value
(Intercept)	0.4895482	0.1456791	3.3604554	0.0007781
modelslogistic regression	0.0000000	0.2060214	0.0000000	1.0000000
modelssingle tree	-0.1255828	0.2046756	-0.6135701	0.5394994
modelsrandom forest	-0.0632055	0.2052927	-0.3078800	0.7581736